

Noah Anderson
July 15th, 2024

Load Libraries

```
In [ ]: from efficientnet_pytorch import EfficientNet
        from torchvision import transforms
        from torchvision.datasets import ImageFolder
        from torch.utils.data.data_loader import DataLoader
        import os
        import torch
        from torch.optim.lr_scheduler import ReduceLROnPlateau
        import multiprocessing
        from torch.utils.data import DataLoader, ConcatDataset, random_split
        import numpy as np
        import torch
        from torchvision.transforms import Compose, RandomRotation, RandomHorizontalFlip, Resize, CenterCrop, ToTensor
        from PIL import Image
        import torch
        from sklearn.metrics import precision_score, recall_score, f1_score
```

Define Transformations and Read Data

```
In [ ]: # Center crop and resize to 224 by 224 and convert to tensor
        data_transform = Compose([
            Resize(224),
            CenterCrop(224),
            ToTensor()
        ])

        # Define data path
        data_dir = "mendeley_data"

        # Read in the data with ImageFolder
        data = ImageFolder(data_dir, transform = data_transform)
```

Shuffling and Re-Splitting Test and Train

```
In [ ]: # Set seed for reproducibility
torch.manual_seed(42) # You can use any integer value as the seed

# Define the sizes for the splits
train_size = int(0.8 * len(data))
val_size = int(0.1 * len(data))
test_size = len(data) - train_size - val_size

# Split the dataset into new train and test sets (e.g., 90% train, 10% test)
trainset, valset, testset = random_split(data, [train_size, val_size, test_size])

# Define batchsize
batch_size = 32

# DataLoaders
train_loader = DataLoader(trainset, batch_size = batch_size, shuffle = True, num_workers = 6, pin_memory =
val_loader = DataLoader(valset, batch_size = batch_size * 2, shuffle = True, num_workers = 6, pin_memory =
test_loader = DataLoader(testset, batch_size = batch_size * 2, num_workers = 6, pin_memory = True)
```

Change Device to GPU

```
In [ ]: # Set device to CPU
device = torch.device("mps") # try runnign mps next
print(f'Using device: {device}')

multiprocessing.set_start_method("fork")
```

Using device: mps

Set Model Features

```
In [ ]: # Load the pre-trained EfficientNet-B0 model
model = EfficientNet.from_pretrained('efficientnet-b0')
```

```

# Modify the last layer to match the number of classes in your dataset
num_classes = len(data.classes)
model._fc = torch.nn.Linear(model._fc.in_features, num_classes)

# Move model to the device (CPU in this case)
model.to(device)

# Define loss function and optimizer
criterion = torch.nn.CrossEntropyLoss()

# Select Adam optimizer
optimizer = torch.optim.Adam(model.parameters(), lr = .001) # Learning rate used in Kaur et. al

# Learning Rate Scheduler
scheduler = ReduceLROnPlateau(optimizer, mode = 'max', factor = 0.3, patience = 2)

# Define accuracy function
def calculate_accuracy(outputs, labels):
    _, preds = torch.max(outputs, 1)
    corrects = (preds == labels).sum().item()
    return corrects

```

Loaded pretrained weights for efficientnet-b0

Train Model

```

In [ ]: # Set initial values for epoch tracking lists
train_losses = [] # Training loss list
val_losses = [] # Validation loss list
train_accuracies = [] # Training accuracy list
val_accuracies = [] # Validation accuracy list
num_epochs = 30 # Define number of epochs
min_valid_loss = np.inf # Initialize minimum validation loss as infinite
max_valid_accuracy = 0 # Initialize maximum validation accuracy as 0

# Epoch training loop
for epoch in range(num_epochs):

```

```

model.train() # Set the model to training mode
train_loss = 0.0 # Initialize training loss
running_corrects = 0

for batch_images, batch_labels in train_loader:
    batch_images, batch_labels = batch_images.to(device), batch_labels.to(device) # Move data to the device

    optimizer.zero_grad() # Zero the gradients
    outputs = model(batch_images) # Forward pass
    loss = criterion(outputs, batch_labels) # Compute loss
    loss.backward() # Backward pass
    optimizer.step() # Update model parameters

    train_loss += loss.item() * batch_images.size(0) # Accumulate loss
    running_corrects += calculate_accuracy(outputs, batch_labels) # Accumulate correct predictions

print("    Epoch training complete, now starting validation")

### Validation
val_loss = 0.0
val_corrects = 0
model.eval()
with torch.no_grad(): # Disable gradient calculation
    for batch_images, batch_labels in val_loader:
        batch_images, batch_labels = batch_images.to(device), batch_labels.to(device) # Move data to the device

        outputs = model(batch_images)
        loss = criterion(outputs, batch_labels)
        val_loss += loss.item() * batch_images.size(0) # Accumulate loss
        val_corrects += calculate_accuracy(outputs, batch_labels) # Accumulate correct predictions

train_loss /= len(train_loader.dataset)
val_loss /= len(val_loader.dataset)
train_acc = running_corrects / len(train_loader.dataset)
val_acc = val_corrects / len(val_loader.dataset)

print(f'Epoch {epoch + 1} \t\t Training Loss: {train_loss:.4f} \t\t Validation Loss: {val_loss:.4f} \t\t')

train_losses.append(train_loss)
val_losses.append(val_loss)
train_accuracies.append(train_acc)
val_accuracies.append(val_acc)

```

```
if min_valid_loss > val_loss:
    print(f'Validation loss decreased ({min_valid_loss:.6f}--->{val_loss:.6f}) \t Saving the model')
    min_valid_loss = val_loss
    torch.save(model.state_dict(), 'saved_models/saved_model.pth')

if val_acc > max_valid_accuracy:
    print(f'Validation accuracy increased ({max_valid_accuracy:.6f}--->{val_acc}) \t Saving the best a
    max_valid_accuracy = val_acc
    torch.save(model.state_dict(), 'saved_models/best_accuracy_model.pth')

scheduler.step(val_loss)
```

```

Epoch training complete, now starting validation
Epoch 1      Training Loss: 0.1921      Validation Loss: 1.8668      Training A
ccuracy: 0.9379      Validation Accuracy: 0.5541
Validation loss decreased(inf-->1.866763)      Saving the model
Validation accuracy increased (0.000000)--->0.5541310541310541      Saving the best accuracy model
Epoch training complete, now starting validation
Epoch 2      Training Loss: 0.0729      Validation Loss: 0.1751      Training A
ccuracy: 0.9753      Validation Accuracy: 0.9345
Validation loss decreased(1.866763--->0.175090)      Saving the model
Validation accuracy increased (0.554131)--->0.9344729344729344      Saving the best accuracy model
Epoch training complete, now starting validation
Epoch 3      Training Loss: 0.0607      Validation Loss: 0.1441      Training A
ccuracy: 0.9810      Validation Accuracy: 0.9387
Validation loss decreased(0.175090--->0.144060)      Saving the model
Validation accuracy increased (0.934473)--->0.9387464387464387      Saving the best accuracy model
Epoch training complete, now starting validation
Epoch 4      Training Loss: 0.0359      Validation Loss: 0.2163      Training A
ccuracy: 0.9897      Validation Accuracy: 0.9416
Validation accuracy increased (0.938746)--->0.9415954415954416      Saving the best accuracy model
Epoch training complete, now starting validation
Epoch 5      Training Loss: 0.0134      Validation Loss: 0.0208      Training A
ccuracy: 0.9959      Validation Accuracy: 0.9972
Validation loss decreased(0.144060--->0.020772)      Saving the model
Validation accuracy increased (0.941595)--->0.9971509971509972      Saving the best accuracy model
Epoch training complete, now starting validation
Epoch 6      Training Loss: 0.0048      Validation Loss: 0.0236      Training A
ccuracy: 0.9989      Validation Accuracy: 0.9900
Epoch training complete, now starting validation
Epoch 7      Training Loss: 0.0045      Validation Loss: 0.0241      Training A
ccuracy: 0.9989      Validation Accuracy: 0.9886
Epoch training complete, now starting validation
Epoch 8      Training Loss: 0.0023      Validation Loss: 0.0195      Training A
ccuracy: 0.9995      Validation Accuracy: 0.9943
Validation loss decreased(0.020772--->0.019507)      Saving the model
Epoch training complete, now starting validation
Epoch 9      Training Loss: 0.0016      Validation Loss: 0.0165      Training A
ccuracy: 0.9996      Validation Accuracy: 0.9929
Validation loss decreased(0.019507--->0.016507)      Saving the model
Epoch training complete, now starting validation
Epoch 10     Training Loss: 0.0017      Validation Loss: 0.0183      Training A
ccuracy: 0.9998      Validation Accuracy: 0.9943
Epoch training complete, now starting validation

```

Epoch 11	Training Loss: 0.0025	Validation Loss: 0.0187	Training A
ccuracy: 0.9995	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 12	Training Loss: 0.0019	Validation Loss: 0.0163	Training A
ccuracy: 0.9996	Validation Accuracy: 0.9929		
Validation loss decreased(0.016507-->0.016280)		Saving the model	
Epoch training complete, now starting validation			
Epoch 13	Training Loss: 0.0032	Validation Loss: 0.0221	Training A
ccuracy: 0.9989	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 14	Training Loss: 0.0007	Validation Loss: 0.0207	Training A
ccuracy: 1.0000	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 15	Training Loss: 0.0011	Validation Loss: 0.0204	Training A
ccuracy: 1.0000	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 16	Training Loss: 0.0009	Validation Loss: 0.0196	Training A
ccuracy: 1.0000	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 17	Training Loss: 0.0012	Validation Loss: 0.0192	Training A
ccuracy: 0.9998	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 18	Training Loss: 0.0013	Validation Loss: 0.0196	Training A
ccuracy: 0.9996	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 19	Training Loss: 0.0014	Validation Loss: 0.0199	Training A
ccuracy: 0.9998	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 20	Training Loss: 0.0009	Validation Loss: 0.0196	Training A
ccuracy: 1.0000	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 21	Training Loss: 0.0014	Validation Loss: 0.0194	Training A
ccuracy: 0.9995	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 22	Training Loss: 0.0009	Validation Loss: 0.0199	Training A
ccuracy: 1.0000	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 23	Training Loss: 0.0016	Validation Loss: 0.0195	Training A
ccuracy: 0.9996	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 24	Training Loss: 0.0011	Validation Loss: 0.0198	Training A
ccuracy: 0.9998	Validation Accuracy: 0.9943		

Epoch training complete, now starting validation			
Epoch 25	Training Loss: 0.0012	Validation Loss: 0.0193	Training A
ccuracy: 0.9995	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 26	Training Loss: 0.0018	Validation Loss: 0.0203	Training A
ccuracy: 0.9996	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 27	Training Loss: 0.0015	Validation Loss: 0.0199	Training A
ccuracy: 0.9996	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 28	Training Loss: 0.0022	Validation Loss: 0.0193	Training A
ccuracy: 0.9993	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 29	Training Loss: 0.0012	Validation Loss: 0.0195	Training A
ccuracy: 0.9996	Validation Accuracy: 0.9943		
Epoch training complete, now starting validation			
Epoch 30	Training Loss: 0.0011	Validation Loss: 0.0203	Training A
ccuracy: 0.9998	Validation Accuracy: 0.9943		

Test Model

Test model after all epochs

```
In [ ]: model.eval() # Set the model to evaluation mode
correct = 0
total = 0
predictions = []
true_labels = []

with torch.no_grad():
    for batch_images, batch_labels in test_loader:
        batch_images, batch_labels = batch_images.to(device), batch_labels.to(device)
        outputs = model(batch_images)
        _, predicted = torch.max(outputs, 1)

        predictions.extend(predicted.cpu().numpy()) # Convert predictions to CPU numpy array
        true_labels.extend(batch_labels.cpu().numpy()) # Convert true labels to CPU numpy array

    total += batch_labels.size(0)
```



```

        correct += (predicted == batch_labels).sum().item()

# Calculate accuracy
accuracy = 100 * correct / total
print(f'Validation Accuracy: {accuracy:.2f}%')

# Calculate precision score
precision = precision_score(true_labels, predictions, average='macro')
print(f'Precision Score: {precision:.4f}')

# Calculate recall score
recall = recall_score(true_labels, predictions, average='macro')
print(f'Recall Score: {recall:.4f}')

# Calculate F1 score
f1 = f1_score(true_labels, predictions, average='macro')
print(f'F1 Score: {f1:.4f}')

```

Validation Accuracy: 99.43%

Precision Score: 0.9948

Recall Score: 0.9942

F1 Score: 0.9945

Test model with the lowest validation loss

```

In [ ]: model.load_state_dict(torch.load('saved_models/saved_model.pth'))

model.eval() # Set the model to evaluation mode
correct = 0
total = 0
predictions = []
true_labels = []

with torch.no_grad():
    for batch_images, batch_labels in test_loader:
        batch_images, batch_labels = batch_images.to(device), batch_labels.to(device)
        outputs = model(batch_images)
        _, predicted = torch.max(outputs, 1)

        predictions.extend(predicted.cpu().numpy()) # Convert predictions to CPU numpy array

```

```

    true_labels.extend(batch_labels.cpu().numpy()) # Convert true labels to CPU numpy array

    total += batch_labels.size(0)
    correct += (predicted == batch_labels).sum().item()

# Calculate accuracy
accuracy = 100 * correct / total
print(f'Validation Accuracy: {accuracy:.2f}%')

# Calculate precision score
precision = precision_score(true_labels, predictions, average='macro')
print(f'Precision Score: {precision:.4f}')

# Calculate recall score
recall = recall_score(true_labels, predictions, average='macro')
print(f'Recall Score: {recall:.4f}')

# Calculate F1 score
f1 = f1_score(true_labels, predictions, average='macro')
print(f'F1 Score: {f1:.4f}')

```

Validation Accuracy: 99.29%

Precision Score: 0.9933

Recall Score: 0.9930

F1 Score: 0.9932

Test model with best validation accuracy

```

In [ ]: model.load_state_dict(torch.load('saved_models/best_accuracy_model.pth'))

model.eval() # Set the model to evaluation mode
correct = 0
total = 0
predictions = []
true_labels = []

with torch.no_grad():
    for batch_images, batch_labels in test_loader:
        batch_images, batch_labels = batch_images.to(device), batch_labels.to(device)
        outputs = model(batch_images)
        _, predicted = torch.max(outputs, 1)

```

```
predictions.extend(predicted.cpu().numpy()) # Convert predictions to CPU numpy array
true_labels.extend(batch_labels.cpu().numpy()) # Convert true labels to CPU numpy array

total += batch_labels.size(0)
correct += (predicted == batch_labels).sum().item()

# Calculate accuracy
accuracy = 100 * correct / total
print(f'Validation Accuracy: {accuracy:.2f}%')

# Calculate precision score
precision = precision_score(true_labels, predictions, average='macro')
print(f'Precision Score: {precision:.4f}')

# Calculate recall score
recall = recall_score(true_labels, predictions, average='macro')
print(f'Recall Score: {recall:.4f}')

# Calculate F1 score
f1 = f1_score(true_labels, predictions, average='macro')
print(f'F1 Score: {f1:.4f}')
```

Validation Accuracy: 99.57%

Precision Score: 0.9960

Recall Score: 0.9957

F1 Score: 0.9959

Plot Training and Validation and Loss and Accuracy

```
In [ ]: import matplotlib.pyplot as plt

# Plotting loss
plt.figure(figsize=(12, 6))

epochs = range(1, len(train_losses) + 1)

plt.subplot(1, 2, 1)
plt.plot(epochs, train_losses, 'b', label='Training loss')
```

```
plt.plot(epochs, val_losses, 'r', label='Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Plotting accuracy
plt.subplot(1, 2, 2)
plt.plot(epochs, train_accuracies, 'b', label='Training accuracy')
plt.plot(epochs, val_accuracies, 'r', label='Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

